

Linear regression

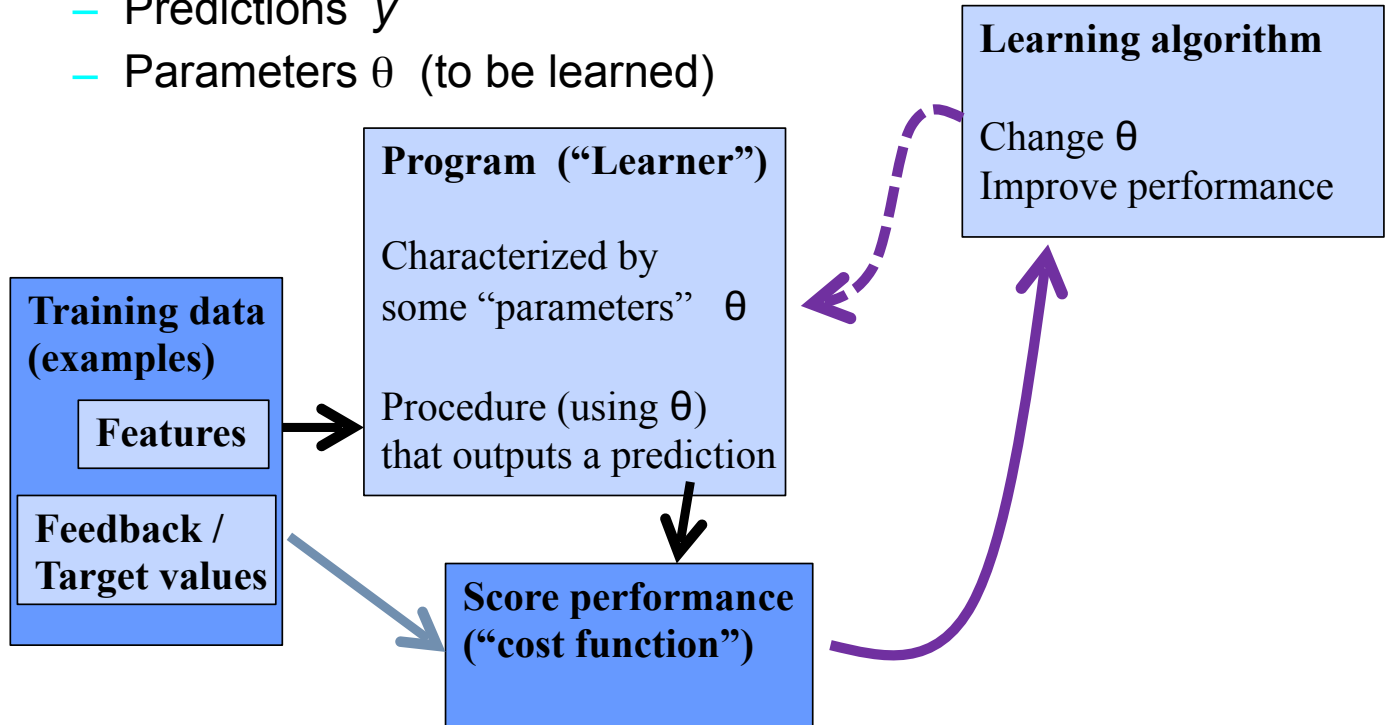
Hung-Hsuan Chen

Adapted from Alexander Ihler (UCI) and Tony R. Martinez (BYU)


Parametric supervised learning

- Notation

- Features x
- Targets y
- Predictions \hat{y}
- Parameters θ (to be learned)



Linear regression





“Predictor”:


$$\hat{y} = \theta_0 + \theta_1 x$$

- Define the form of the function $f(x)$ explicitly
 - i.e., $\hat{y} = \theta_0 + \theta_1 x$ in this case
 - If you believe that x and y have a **non-linear** relationship, you should assume a non-linear $f(x)$
- Find a good $f(x)$ within that family
 - i.e. find good θ_0 and θ_1 such that $\hat{y} \approx y$ in this case

Quiz: which one should be an error measurement?



Measuring error



Simple linear regression

- For now, assume just one (input) independent variable x , and one (output) dependent variable y
 - Multiple linear regression assumes an input **vector** \mathbf{x}
- We will "fit" the points with a line
 - In a high dimensional space, we will fit the points with a **hyper-plane**
- Which line should we use?
 - Choose an objective function
 - Typically, we choose sum squared error (SSE) (or its variants)
 - $\frac{1}{n} \sum (\hat{y}_i - y_i)^2 = \frac{1}{n} \sum (\text{residual}_i)^2$
 - Choices with the same result: $\frac{1}{2n} \sum (\hat{y}_i - y_i)^2$ or $\frac{1}{2} \sum (\hat{y}_i - y_i)^2$
 - You may choose other objective functions
 - More to come in future lectures

why $\frac{1}{2}$? 微分较
好看

How to “learn” the parameters?

- For the 2- d problem there are coefficients for the bias and the independent variable (y-intercept and slope)

- $\hat{y} = \theta_0 + \theta_1 x$


- The value of the coefficients which minimize the objective function:

- $$\theta_1 = \frac{n \sum x_i y_i - \sum x_i \sum y_i}{n \sum x_i^2 - (\sum x_i)^2}$$

- $$\theta_0 = \frac{\sum y_i - \theta_1 \sum x_i}{n}$$

- n : number of training instances

Visualizing θ_0 and θ_1



How to derive the value of the coefficients?

- $J(\theta_0, \theta_1) = \frac{1}{2n} \sum (\hat{y}_i - y_i)^2$
 $= \frac{1}{2n} \sum ((\theta_0 + \theta_1 x_i) - y_i)^2,$



- Set

- $\frac{\partial J(\theta_0, \theta_1)}{\partial \theta_0} = 0$ and

- $\frac{\partial J(\theta_0, \theta_1)}{\partial \theta_1} = 0$

to solve θ_0 and θ_1

More dimensions?



$$\hat{y}(\mathbf{x}_i) = \theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2} + \dots + \theta_d x_{i,d} = \boldsymbol{\theta}^T \mathbf{x}_i, \text{ where}$$

$$\mathbf{x}_i = [1, x_{i,1}, x_{i,2}, \dots, x_{i,d}]^T$$

$$\boldsymbol{\theta} = [\theta_0, \theta_1, \theta_2, \dots, \theta_d]^T$$

d : # of features, $x_{i,j}$: the i th training instance's j th feature

Multiple linear regression

- n : the number of training instances
- d : the number of features
- Training instances:

$$\mathbf{X} = \begin{bmatrix} x_{1,1} & \cdots & x_{1,d} \\ \vdots & \ddots & \vdots \\ x_{n,1} & \cdots & x_{n,d} \end{bmatrix}, \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}$$

— (We assume no coefficient parameter here)

- Find $\boldsymbol{\theta} = \begin{bmatrix} \theta_1 \\ \vdots \\ \theta_d \end{bmatrix}$ such that $(\hat{\mathbf{y}} - \mathbf{y})^T (\hat{\mathbf{y}} - \mathbf{y})$ is minimized,

where

$$\hat{\mathbf{y}} = \mathbf{X}\boldsymbol{\theta}$$

- The solution is $\boldsymbol{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

How to derive the value of the coefficient vector?

- $J(\theta) = \frac{1}{2}(\hat{\mathbf{y}} - \mathbf{y})^T (\hat{\mathbf{y}} - \mathbf{y})$
 $= \frac{1}{2}(\mathbf{X}\theta - \mathbf{y})^T (\mathbf{X}\theta - \mathbf{y})$
- Set
 - $\nabla J(\theta) = \mathbf{0}$to solve θ



Another way to find θ for the linear regression problem

- Goal: find θ_0 and θ_1 to minimize

$$\begin{aligned} J(\theta_0, \theta_1) &= \frac{1}{2n} \sum (\hat{y}_i - y_i)^2 \\ &= \frac{1}{2n} \sum ((\theta_0 + \theta_1 x_i) - y_i)^2 \end{aligned}$$

- Procedure
 1. Start with $\theta_0 = r_0$ (a random number) and $\theta_1 = r_1$ (another random number)
 2. Slightly move θ_0 and θ_1 to reduce $J(\theta_0, \theta_1)$
 3. Keep doing step 2 until converged
- Question: how to move θ_0 and θ_1 ?


Plotting the objective function (3D and contour plots)



Keep moving "downward" to reach the minimum

Finding good parameters

- Want to find parameters which minimize our error...
- Think of a cost “surface”: error residual for that θ ...



Gradient descent

$$\theta_0^{(k)} = 5$$

$$\theta_1^{(k)} = -4$$

$$\hat{y} = \theta_0 + \theta_1 x$$

$$J = \left(y - (\theta_0 + \theta_1 x) \right)^2$$

$$\frac{\partial J}{\partial \theta_0} = 2(\downarrow)(-1)$$

$$\frac{\partial J}{\partial \theta_1} = 2(\downarrow)(-x)$$

• Procedure

1. Start with random values

- $\theta = \theta^{(0)} = (\theta_0^{(0)}, \theta_1^{(0)}, \dots, \theta_d^{(0)})$

2. Slightly move $\theta_0, \dots, \theta_d$ to reduce $J(\theta)$

- $$\theta_i^{(k+1)} = \theta_i^{(k)} - \alpha \frac{\partial J(\theta)}{\partial \theta_i} \Big|_{\theta = \theta^{(k)}}$$

α is a small positive number

- $k = k + 1$

3. Keep doing step 2 until convergence

$$\frac{\partial}{\partial \theta_i} J(\theta) = \frac{1}{n} \sum_j (\hat{y}_j - y_j) x_{ji},$$

$\frac{1}{n}$ is usually ignored


Why gradient descent work?

Repeat until converge {



for (i=1,2, ..., d){

$$\theta_i := \theta_i - \alpha \frac{\partial}{\partial \theta_i} J(\theta) = \theta_i - \alpha \sum_j (\hat{y}_j - y_j) x_{ji}$$

}




Gradient descent (one feature and two features)



Overly large learning rate may not lead to converge

- $$\theta_i^{(k+1)} := \theta_i^{(k)} - \alpha \frac{\partial}{\partial \theta_i} J(\theta^{(k)})$$



- α : learning rate
- Often $\alpha \in [0.001, 1]$
- Shrink α as k becomes larger

We now have two tools to solve linear regression problem

- Close form solution

$$\theta = (X^T X)^{-1} X^T y$$

- Gradient descent

```
Repeat until converge {  
  for (i=1,2, ..., d){  
     $\theta_i = \theta_i - \alpha \sum_j (\hat{y}_j - y_j) x_{ji}$   
  }  
}
```

Which one should we use?

Consider using close form solution

- $\theta = (X^T X)^{-1} X^T y$
- Required space:
 - Even if X is sparse, $(X^T X)^{-1}$ may not be sparse
 - If $n=1M$, $d=100K$, $(X^T X)^{-1}$'s dimension is $(100k)^2 = 10^{10}$
 - If 8 bytes per entry, storing $(X^T X)^{-1}$ requires 80GB \rightarrow typically infeasible on a single machine
- Computation time:
 - Computing matrix multiplication $X^T X$ takes $O(nd^2)$
 - Let $(X^T X) = P$, computing P^{-1} takes $O(d^{2.373})$ to $O(d^3)$
 - Depending on the inversion algorithm
 - Let $(X^T X)^{-1} = Q$, computing QX^T takes $O(nd^2)$
 - Let $(X^T X)^{-1} X^T = R$, computing Ry takes $O(nd)$
 - Total: $O(nd^2 + d^{2.373} + nd^2)$
- If d is small, close form solution is probably acceptable

n: number of

Instance

d: number of

Features

Consider using Gradient Descent

```
Repeat until converge {  
  for (j=1, 2, ..., n){  
     $z_j = (\boldsymbol{\theta}^T \mathbf{x}_j - y_j)$   
  }  
  for (i=1, 2, ..., d){  
     $\theta_i = \theta_i - \alpha \sum_j z_j x_{ji}$   
  }  
}
```

- Required space:
 - $\mathbf{X} = [x_{ij}]$, \mathbf{y} , $\boldsymbol{\theta}$, and \mathbf{z}
 - If $n=1\text{M}$, $d=100\text{K}$, and \mathbf{X} is sparse
 - If 8 bytes per entry, we need: $(\text{nnz}(\mathbf{X}) + 1\text{M} + 100\text{K} + 1\text{M}) * 8\text{bytes}$
 - Much more efficient than 80GB
- Computation time:
 - $z_j = (\boldsymbol{\theta}^T \mathbf{x}_j - y_j)$ for all j :
 $O(nd)$
 - $\theta_i := \theta_i - \alpha \sum_j z_j x_{ji}$ for all i :
 $O(nd)$
 - Outer loop: Assume T iterations
 - Total: $O(Tnd)$
 - Usually more efficient than $O(nd^2 + d^{2.373} + nd^2)$

Close form vs gradient descent

- If the number of features is small, close form solution is probably acceptable
- However, if the number of features is large, using gradient descent is more efficient in both memory usage and computing time
- Moreover, gradient descent is capable of solving more complex optimization problem

— In many cases, $\frac{\partial J(\theta)}{\partial \theta} = \mathbf{0}$ has no closed-form solution

- But we can still apply gradient descent 😊

— More to come in future lectures

Quiz

- If the features are all **categorical**
 - Can you apply linear regression?
 - Can you apply decision tree?
 - Can you apply knn?
- If all features are **numerical**, which one runs faster during “test” phase?
 - Linear regression? Decision tree? KNN?

Lg Little faster than Dt

When to stop updating?

- Possible stopping criteria:
 - Improvement drops (e.g., < 0)
 - Reached small error
 - Achieved predefined # of iterations
 - No time to train anymore


Fitting non-linear data to linear model

- $\hat{y}(\mathbf{x}_i) = \theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2} + \dots + \theta_d x_{i,d}$
 - Linear model
- We may generate the higher degree terms as the new features
 - $\hat{y}(\mathbf{x}_i) = (\theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2} + \dots + \theta_d x_{i,d}) +$
 $(\theta_{1,2} x_{i,1} x_{i,2} + \theta_{1,3} x_{i,1} x_{i,3} + \dots + \theta_{d-1,d} x_{i,d-1} x_{i,d}) +$
 $(\theta_{1,1} x_{i,1}^2 + \theta_{2,2} x_{i,2}^2 + \dots + \theta_{d,d} x_{i,d}^2)$


Overfitting

- Overfitting occurs when a model captures idiosyncrasies of the input data, rather than generalizing
 - Too many parameters relative to the amount of training data
- For example, an order- N polynomial can perfectly fit to $N+1$ data points

Target: $\sin(2\pi x)$ + noise





$$\hat{y}_i = \theta_0$$



$$\hat{y}_i = \theta_0 + \theta_1 x$$

Target: $\sin(2\pi x) + \text{noise}$ (overfitting)

$$\hat{y}_i = \theta_0 + \sum_{i=1}^9 \theta_i x^i$$



$$\hat{y}_i = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3$$

Observation

- In the linear regression model, overfitting is characterized by large parameters

	$d = 0$	$d = 1$	$d = 3$	$d = 9$
	0.19	0.82	0.31	0.35
		-1.27	7.99	232.37
			-25.43	-5321.83
			17.37	48568.31
				231639.30
				640042.26
				-1061800.52
				1042400.18
				-557682.99
				125201.43

Regularization to avoid overfitting

- Introduce a penalty term for the size of the weights
- Un-regularized regression (the original objective function)

$$- J(\boldsymbol{\theta}) = \frac{1}{2} \sum (\hat{y}_i - y_i)^2 = \frac{1}{2} \sum (\boldsymbol{\theta}^T \mathbf{x}_i - y_i)^2$$

$$- \boldsymbol{\theta} := \operatorname{argmin}_{\boldsymbol{\theta}} \frac{1}{2} \sum (\hat{y}_i - y_i)^2$$

- Regularized regression (enforce the solution to have low L2-norm of $\boldsymbol{\theta}$)

$$- \boldsymbol{\theta} := \operatorname{argmin}_{\boldsymbol{\theta}} \frac{1}{2} \sum (\hat{y}_i - y_i)^2 \text{ such that } \|\boldsymbol{\theta}\|^2 \leq K$$

$$\theta_1^2 + \theta_2^2 + \dots + \theta_n^2$$

Regularization to avoid overfitting

- Original objective function: minimizing the training error

- $J(\boldsymbol{\theta}) = \frac{1}{2} \sum (\hat{y}_i - y_i)^2 = \frac{1}{2} \sum (\boldsymbol{\theta}^T \mathbf{x}_i - y_i)^2$

- New target

- $\boldsymbol{\theta} := \operatorname{argmin}_{\boldsymbol{\theta}} \frac{1}{2} \sum (\hat{y}_i - y_i)^2$ such that $\|\boldsymbol{\theta}\|^2 \leq K$

- This is equivalent to the following problem with some λ

- $\boldsymbol{\theta} := \operatorname{argmin}_{\boldsymbol{\theta}} \left[\frac{1}{2} \sum (\hat{y}_i - y_i)^2 + \underbrace{\frac{\lambda}{2} \|\boldsymbol{\theta}\|^2}_{\text{hyperparameter}} \right]$

- New objective function: minimizing training error and the L2-norm of $\boldsymbol{\theta}$ simultaneously

- Solution: $\boldsymbol{\theta} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$

* λ 越大
得到的 $\boldsymbol{\theta}$ 小

Regularized linear regression

- Regularized linear regression

- $\theta := \operatorname{argmin}_{\theta} \left[\frac{1}{2} \sum (\hat{y}_i - y_i)^2 + R(\theta) \right]$

- $R(\theta)$: regularization

- L2-regularization (**Ridge regression**)

- $\theta := \operatorname{argmin}_{\theta} \left[\frac{1}{2} \sum (\hat{y}_i - y_i)^2 + \lambda \|\theta\|^2 \right]$
 $\rightarrow \theta_1^2 + \theta_2^2 \dots$

- L1-regularization (**Lasso**)

- $\theta := \operatorname{argmin}_{\theta} \left[\frac{1}{2} \sum (\hat{y}_i - y_i)^2 + \lambda \|\theta\|_1 \right]$
 $\rightarrow |\theta_1| + |\theta_2| + \dots$

Combining Ridge and Lasso

- Elastic net regularization

- $\boldsymbol{\theta} := \operatorname{argmin}_{\boldsymbol{\theta}} \left[\sum \left(\hat{y}_i - y_i \right)^2 + \lambda_1 \|\boldsymbol{\theta}\|_1 + \lambda_2 \|\boldsymbol{\theta}\|^2 \right]$

- When $\lambda_1 = 0$ and $\lambda_2 > 0 \rightarrow$ Ridge regression

- When $\lambda_1 > 0$ and $\lambda_2 = 0 \rightarrow$ Lasso

How to select λ ?

- Split the data into training and test datasets
 - Based on the training data, train the models by different λ 's
 - Based on the test data, calculate the test performance of all the models (of different λ 's)
 - Select the λ with the best test performance
- Cross validation
 - More to come in future lectures



Properties of Ridge and Lasso

- Ridge

每个 feature 皆有 effect

- Good if many features have small/medium sized effects


- Lasso

只有特定 feature 有 effect

- Good if only a few features with a medium/large effect

Geometry of Ridge, Lasso, and Elastic net


Elastic net: $\lambda_1 = \lambda_2 = 0.5$



ex: $|\theta_1|^{0.01} + |\theta_2|^{0.01} \leq k$

Why Lasso zeros coefficients

解釋 1



LASSO

$|\theta_1| + |\theta_2| \leq k$

RIDGE REGRESSION

$\theta_1^2 + \theta_2^2 \leq k$

Why Lasso zeros coefficients? – a numerical explanation

解釋 2

- Ridge

– $\boldsymbol{\theta} := \operatorname{argmin}_{\boldsymbol{\theta}} \left[\frac{1}{2} \sum \left(\hat{y}_i - y_i \right)^2 + \lambda \|\boldsymbol{\theta}\|^2 \right]$

- Changing θ_j from 2 to 1 reduces the cost by $\lambda(22 - 12) = 3\lambda$
- Changing θ_j from 1 to 0 reduces the cost by $\lambda(12 - 02) = \lambda$


- Lasso

– $\boldsymbol{\theta} := \operatorname{argmin}_{\boldsymbol{\theta}} \left[\frac{1}{2} \sum \left(\hat{y}_i - y_i \right)^2 + \lambda \|\boldsymbol{\theta}\|_1 \right]$

- Changing θ_j from 2 to 1 reduces the cost by $\lambda(2 - 1) = \lambda$
- Changing θ_j from 1 to 0 reduces the cost by $\lambda(1 - 0) = \lambda$

→ Ridge tends to shrink large coefficients to smaller ones, but not to shrink small coefficients to zero

Generalize the regularization term



- $\theta := \operatorname{argmin}_{\theta} \left[\frac{1}{2} \sum \left(\hat{y}_i - y_i \right)^2 + \lambda \|\theta\|^q \right]$

$$q \geq 0$$

- $q=1$: Lasso
- $q=2$: Ridge regression
- A smaller q tends to shrink the coefficients

Gradient descent for ridge regression

- Gradient descent (GD)

- $\boldsymbol{\theta}^{(k+1)} := \boldsymbol{\theta}^{(k)} - \alpha \nabla J(\boldsymbol{\theta}^{(k)})^T$

- Ridge regression

- $J(\boldsymbol{\theta}) = \frac{1}{2}(\hat{\mathbf{y}} - \mathbf{y})^T (\hat{\mathbf{y}} - \mathbf{y}) + \frac{\lambda}{2} \|\boldsymbol{\theta}\|^2$

$$= \frac{1}{2}(\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y}) + \frac{\lambda}{2} \boldsymbol{\theta}^T \boldsymbol{\theta}$$

- $\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) = (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})^T \mathbf{X} + \lambda \boldsymbol{\theta}^T$


- $\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta})^T = \mathbf{X}^T (\mathbf{X}\boldsymbol{\theta} - \mathbf{y}) + \lambda \boldsymbol{\theta}$

Gradient descent for ridge regression (pseudo code)

```
1   $k := 0$   
2  Initialize  $\boldsymbol{\theta}^{(k)}$   
3  while not converge:  
4       $\mathbf{g} := \mathbf{X}^T(\mathbf{X}\boldsymbol{\theta}^{(k)} - \mathbf{y}) + \lambda\boldsymbol{\theta}^{(k)}$   
5       $\boldsymbol{\theta}^{(k+1)} := \boldsymbol{\theta}^{(k)} - \alpha\mathbf{g}$   
6       $k := k+1$ 
```


Convex function

- Convex function
 - Local minimum is global minimum
- The least-square linear regression objective and the Ridge regression objective are both convex functions
 - Gradient descent will find the unique minimum (or very close, depending on the step size α)



Effect of the step size α

- Small α : Long computation time
- Large α : may not reach the minimum




- A common practice
 - Gradually shrink the value of the step size

Stochastic gradient descent (SGD) – motivation

- In GD, the gradient is computed by
 - $\mathbf{g} := \mathbf{X}^T(\mathbf{X}\boldsymbol{\theta}^{(k)} - \mathbf{y}) + \lambda\boldsymbol{\theta}^{(k)}$
 - The entire training set is examined at each step
 - Very slow when the # of training instances is large!

```
1   $k := 0$ 
2  Initialize  $\boldsymbol{\theta}^{(k)}$ 
3  while not converge:
4       $\mathbf{g} := \mathbf{X}^T(\mathbf{X}\boldsymbol{\theta}^{(k)} - \mathbf{y}) + \lambda\boldsymbol{\theta}^{(k)}$ 
5       $\boldsymbol{\theta}^{(k+1)} := \boldsymbol{\theta}^{(k)} - \alpha\mathbf{g}$ 
6       $k := k+1$ 
```



Stochastic gradient descent

- Optimize **one example** at a time

1 $k := 0$

2 Initialize $\boldsymbol{\theta}^{(k)}$

Only one training instance is
examined at each step

3 while not converge:

4 $\mathbf{g} := \mathbf{x}_i^T (\mathbf{x}_i^T \boldsymbol{\theta}^{(k)} - \mathbf{y}_i) + \lambda \boldsymbol{\theta}^{(k)}$

5 $\boldsymbol{\theta}^{(k+1)} := \boldsymbol{\theta}^{(k)} - \alpha \mathbf{g}$

6 $k := k+1$



7 Get next data instance i

- Why does SGD work?

$$E\left(\nabla J_i(\boldsymbol{\theta})\right) = \frac{1}{n} \sum \nabla J_i(\boldsymbol{\theta}) = \nabla J(\boldsymbol{\theta})$$

Gradient descent (GD) vs stochastic gradient descent (SGD)

- # steps
 - GD: fewer steps
 - SGD: more steps
- Computation of each step
 - GD: look through all the training instances
 - SGD: look only one training instance



Pros and cons of SGD

- Pros
 - When the training data is large with some (near) redundant instances, SGD is usually much faster to converge than GD
 - Supports online learning
- Cons
 - Tends to bouncing around the minimum

Mini-batch gradient descent

- Optimize ***few examples*** at a time

1 $k := 0$

2 Initialize $\theta^{(k)}$

Training instances $(i, i+1, \dots, j)$
are examined

3 while not converge:

4 $g := \mathbf{x}_{i:j}^T \left(\mathbf{x}_{i:j} \theta^{(k)} - \mathbf{y}_{i:j} \right) + \lambda \theta^{(k)}$

5 $\theta^{(k+1)} := \theta^{(k)} - \alpha g$

6 $k := k+1$

7 Get next batch $(i, i+1, \dots, j)$

Gradient descent/stochastic gradient descent/mini-batch gradient descent

- All of them iteratively update the parameters such that the target function gradually becomes smaller
- If we have n training instances
 - (Batch) gradient descent: every parameter update requires seeing all training instances once
 - Stochastic gradient descent: every parameter update requires seeing one of n training instances
 - Mini-batch: every parameter update requires seeing b training instances (if batch size = b)

Locally weighted linear regression

- Linear regression

- Method

- Assume $\hat{y} = \theta_0 + \theta_1 x$

- Find θ_0 and θ_1 to minimize $\sum_i (y_i - \hat{y}_i)^2$

- Given a new instance x^* , we predict the corresponding y^* as $\theta_0 + \theta_1 x^*$

- Once training finished, we get θ_0 and θ_1 , and training data can be removed

- However, can only model linear relationship

- Locally weighted linear regression

- Model the "nonlinear" data by weighting errors differently

Locally weighted linear regression

- Method

- Given a new instance x^* , we define the weight of each training error by

- $$w_i = \exp\left(\frac{(x_i - x^*)^2}{2\sigma^2}\right)$$


- σ is a hyperparameter; we set to 1 for simplicity
 - If $x_i \approx x^*$, $w_i \approx 1$
 - If x_i and x^* is far, $w_i \approx 0$

- Find θ_0 and θ_1 to minimize

- $$\sum_i w_i \left(y_i - \hat{y}_i\right)^2, \text{ where } \hat{y}_i = \theta_0 + \theta_1 x_i$$

- Adv: can fit nonlinear curve
- Disadv: need to “train” the model for every prediction


LR vs locally weighted LR



Newton's method for optimization (1/2)

- Let $\ell(\theta) = \sum_i (y_i - \hat{y}_i)^2$, we want to find θ to minimize $\ell(\theta)$

— The same as find θ such that $\ell'(\theta) = 0$




Newton's method for optimization (2/2)

- Initial guess: $\theta^{(0)}$
- The slope of the tangent line at $(\theta^{(0)}, \ell'(\theta^{(0)}))$ is $\ell''(\theta^{(0)})$

— i.e., $\ell''(\theta^{(0)}) = \frac{a}{b} = \frac{\ell'(\theta^{(0)})}{\theta^{(0)} - \theta^{(1)}}$

- So, $\theta^{(1)} = \theta^{(0)} - \frac{\ell'(\theta^{(0)})}{\ell''(\theta^{(0)})}$



Newton's method vs SGD

- Newton's method usually requires fewer steps
- However, the cost of each step is usually large
 - If θ is a scalar, using Newton's method requires to compute $f''(\theta)$
 - If θ is a vector, using Newton's method requires to compute the Hessian matrix

Evaluation (regression)

Evaluating regression models

- Mean Square Error (MSE)

- $MSE(\mathbf{y}, \hat{\mathbf{y}}) = \frac{\sum (y_i - \hat{y}_i)^2}{n}$

- Similar metric: Root Mean Squared Error (RMSE)

- Criticism:

- Can only be compared between models whose errors are measured in the same unit
 - Tend to be influenced by extreme values

- Mean Absolute Error (MAE)

- $MAE(\mathbf{y}, \hat{\mathbf{y}}) = \frac{\sum |y_i - \hat{y}_i|}{n}$

- Usually smaller than MSE

- Criticism:

- Can only be compared between models whose errors are measured in the same unit

Evaluating regression models

- Median Absolute Error (MedAE)

- $MedAE(\mathbf{y}, \hat{\mathbf{y}}) = \text{median}(|y_1 - \hat{y}_1|, |y_2 - \hat{y}_2|, \dots, |y_n - \hat{y}_n|)$
- Robust to the extreme values
- Criticism:
 - Can only be compared between models whose errors are measured in the same unit

- R^2 score (coefficient of determination)

- $$R^2(\mathbf{y}, \hat{\mathbf{y}}) = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$$
- The scale of the model is more intuitive
 - Best result: 1
 - Can be arbitrarily worse (very negative)

A closer look of R^2 score (coefficient of determination)

- Mean model

- A naïve regression function – a constant function that always outputs \bar{y} the mean of the target variable in the training dataset

- $$R^2(\mathbf{y}, \hat{\mathbf{y}}) = 1 - \frac{\frac{1}{n} \sum (y_i - \hat{y}_i)^2}{\frac{1}{n} \sum (y_i - \bar{y})^2} = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}$$

- $\sum (y_i - \hat{y}_i)^2$: sum of square error (SSE)

- $\sum (y_i - \bar{y})^2$: sum of square total (SST)

- i.e., SSE of the mean model

- Max value: 1

- When SSE=0 (all predictions are correct)

- Min value: $-\infty$

- SSE can be arbitrarily worse

- $$R^2(\mathbf{y}, \hat{\mathbf{y}}) = 0$$

- The performance of the prediction is the same as the mean model

R (correlation coefficient) vs R^2 (coefficient of determination)

- $\left[R(\mathbf{y}, \hat{\mathbf{y}}) \right]^2$ equals $R^2(\mathbf{y}, \hat{\mathbf{y}})$ when
 - Applying the linear regression model, i.e., $\hat{\mathbf{y}} = \mathbf{X}\boldsymbol{\theta}$, where $\boldsymbol{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$, and
 - Evaluating $R(\mathbf{y}, \hat{\mathbf{y}})$ and $R^2(\mathbf{y}, \hat{\mathbf{y}})$ on the training data
- In this case, $0 \leq R^2(\mathbf{y}, \hat{\mathbf{y}}) \leq 1$

Conclusion (1/3)

- Linear regression

- $\hat{y}_i = \theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2} + \dots + \theta_d x_{i,d}$

- Find good $\theta_0, \theta_1, \dots, \theta_d$ such that $\hat{y}_i \approx y_i \quad \forall i$

- Commonly used distance between \hat{y}_i and y_i : RSS

- Linear regression with constraints on θ

- Prevent overfitting

- Limiting θ by L2-norm: Ridge regression

- Limiting θ by L1-norm: Lasso

- Feature selection

- Limiting θ by both L1-norm and L2-norm: Elastic net

Conclusion (2/3)

- Convex function
 - Local minimum equals global minimum
 - Iteratively adjust $\boldsymbol{\theta}$ to reduce $J(\boldsymbol{\theta})$
 - $\boldsymbol{\theta}^{(k+1)} := \boldsymbol{\theta}^{(k)} - \alpha \mathbf{g}$
- For ridge regression
 - Gradient descent
 - $\mathbf{g} := \mathbf{X}^T(\mathbf{X}\boldsymbol{\theta}^{(k)} - \mathbf{y}) + \lambda\boldsymbol{\theta}^{(k)}$
 - Stochastic gradient descent
 - $\mathbf{g} := \mathbf{x}_i^T(\mathbf{x}_i^T\boldsymbol{\theta}^{(k)} - \mathbf{y}_i) + \lambda\boldsymbol{\theta}^{(k)}$
 - Mini-batch gradient descent
 - $\mathbf{g} := \mathbf{x}_{i:j}^T(\mathbf{x}_{i:j}\boldsymbol{\theta}^{(k)} - \mathbf{y}_{i:j}) + \lambda\boldsymbol{\theta}^{(k)}$

Conclusion (3/3)

- Common metrics to evaluate regression models
 - Mean square error
 - R^2 score

Quiz

- If the training data contain millions of features, should we use gradient descent or closed form solution for training?
- Can gradient descent get stuck in a local minimum (which is not a global minimum) when training a linear regression model?
- Why would one use Ridge regression instead of plain linear regression?
- If the R^2 score on test data is negative, what does this mean?
- If the R^2 score on training data is negative, what does this mean?