

Differentiating regularization weights for recommender systems

Hung-Hsuan Chen

11/24/20


1

Cold start

- New items or the long tail items are rarely (or never) rated/viewed/purchased
- New users or less active users rarely rate/view/purchase items
- We have limited information on these users and items
- Predictions are usually inaccurate, since we have limited (or no) information

2

Item's received rating counts in training data vs test RMSE



(a) The Epinions dataset

3

Pearson's correlation coefficient (PCC) on different datasets

Dataset	PCC	p -value
Epinions	-0.45	1.5×10^{-16} (***)
MovieLens-100K	-0.70	1.8×10^{-37} (***)
FilmTrust	-0.73	2.0×10^{-14} (***)
Yahoo! Movies	-0.12	0.04 (*)
AMI	-0.36	3.6×10^{-10} (***)

4

Key idea

- We have much information of the popular items and the active users
 - We probably should trust the data, since data size is large (less likely to overfit)
 - We should **assign lower constraints** when learning the vectors of these users and items
- We have little information of the long tail items and the less active users
 - We probably should make conservative predictions; trust the data too much may cause overfitting
 - We should **assign larger constraints** when learning the vectors of these users and items

5

Original loss function

- $$L(\Theta) = \frac{1}{2} \sum_{\forall (i,j) \in \tilde{k}} (r_{ij} - \hat{r}_{ij})^2 + \frac{\lambda}{2} \|\Theta\|^2$$

$$= \frac{1}{2} \sum_{\forall (i,j) \in \tilde{k}} (r_{ij} - \hat{r}_{ij})^2 + \frac{\lambda}{2} \left\{ \sum_{\forall i \in K_U} \|p_i\|^2 + \sum_{\forall i \in K_I} \|q_i\|^2 + \sum_{\forall i \in K_U} \|b_i\|^2 + \sum_{\forall i \in K_I} \|c_i\|^2 \right\}$$
- \hat{r}_{ij} can be predicted by Simon Funk's SVD, Factorization Machines, FPMC, etc.
- A universal regularization weight λ is assigned to all parameters

6

New loss function

- $$L(\Theta) = \frac{1}{2} \sum_{\forall (i,j) \in \tilde{k}} (r_{ij} - \hat{r}_{ij})^2 + \frac{\lambda}{2} \left\{ \sum_{\forall i \in K_U} \frac{1}{f(R^{(U)}(i))} \|p_i\|^2 + \sum_{\forall i \in K_I} \frac{1}{f(R^{(I)}(j))} \|q_i\|^2 + \sum_{\forall i \in K_U} \frac{1}{f(R^{(U)}(i))} \|b_i\|^2 + \sum_{\forall i \in K_I} \frac{1}{f(R^{(I)}(j))} \|c_i\|^2 \right\}$$
 - $R^{(U)}(i)$: the number of items rated by user i
 - $R^{(I)}(j)$: the number of users who rated item j
 - $f(x)$ needs to be a positive and monotonically increasing function

7

Regularization terms

- $$\frac{\lambda}{2} \left\{ \sum_{\forall i \in K_U} \frac{1}{f(R^{(U)}(i))} \|p_i\|^2 + \sum_{\forall i \in K_I} \frac{1}{f(R^{(I)}(j))} \|q_i\|^2 + \sum_{\forall i \in K_U} \frac{1}{f(R^{(U)}(i))} \|b_i\|^2 + \sum_{\forall i \in K_I} \frac{1}{f(R^{(I)}(j))} \|c_i\|^2 \right\}$$
- When $R^{(U)}(i)$ is large, we have smaller constraints on p_i and b_i
- When $R^{(I)}(j)$ is large, we have smaller constraints on q_j and c_j

8

Proposed regularization functions

- Linear: $f(x) = ax + b$ ($a > 0, b \geq 0$)
- Logarithm: $f(x) = \log(x + c)$ ($c > 0$)
- Square root: $f(x) = \sqrt{x + c}$ ($c \geq 0$)
- Support of functions: $x \geq 1$

9

Statistics of the benchmark datasets

Dataset	# users	# items	# ratings	Density	Rating scale
Epinions	40,163	139,738	664,824	0.0118%	[1, 2, 3, 4, 5]
MovieLens-100K	943	1,682	100,000	6.3047%	[1, 2, 3, 4, 5]
FilmTrust	1,508	2,071	35,497	1.1366%	[0.5, 1, 1.5, ..., 4]
Yahoo! Movies	7,642	11,916	221,367	0.2431%	[1, 2, ..., 13]
AMI	339,231	83,046	500,176	0.0018%	[1, 2, 3, 4, 5]

10

RMSE scores of SVD and SVD with regularization weights on the test datasets

Dataset	SVD	linear-reg	sqrt-reg	log-reg	improve ratio range
Epinions	1.1997	1.0538	1.0538	1.0538	12.16%
MovieLens-100K	0.9423	0.9422	0.9422	0.9422	0.01%
FilmTrust	0.8465	0.8194	0.8194	0.8223	2.86% to 3.20%
Yahoo! Movies	3.0799	2.9892	3.0129	3.0127	2.18% to 2.94%
AMI	1.1450	1.1405	1.1405	1.1405	0.39%

11

RMSE scores of SVD++ and SVD++ with regularization weights on the test datasets

Dataset	SVD++	linear-reg	sqrt-reg	log-reg	improve ratio range
Epinions	1.0628	1.0538	1.0566	1.0538	0.58% to 0.85%
MovieLens-100K	0.9423	0.9422	0.9422	0.9423	0.00% to 0.01%
FilmTrust	0.8199	0.8199	0.8197	0.8194	0.00% to 0.06%
Yahoo! Movies	3.0152	3.0127	3.0126	3.0128	0.08% to 0.09%
AMI	1.1446	1.1407	1.1429	1.1408	0.15% to 0.34%

12

RMSE scores of NMF and NMF with regularization weights on the test datasets

Dataset	NMF	linear-reg	sqr-reg	log-reg	improve ratio range
Epinions	1.0709	1.0582	1.0582	1.0583	1.18% to 1.19%
MovieLens-100K	0.9621	0.9416	0.9416	0.9413	2.13% to 2.16%
FilmTrust	0.9295	0.8206	0.8206	0.8206	11.72%
Yahoo! Movies	3.6058	3.0140	3.0141	3.0141	16.41%
AMI	1.1431	1.1433	1.1432	1.1432	-0.01% to -0.02%

13

Various metrics of SVD and SVD + regularization weights on the Epinions dataset

Metrics	SVD	linear-reg	sqr-reg	log-reg	improve ratio range
RMSE	1.1997	1.0538	1.0538	1.0538	12.16%
MAE	0.9139	0.8190	0.8190	0.8190	10.38%
R ²	0.0067	0.2337	0.2337	0.2337	3388.06%

14

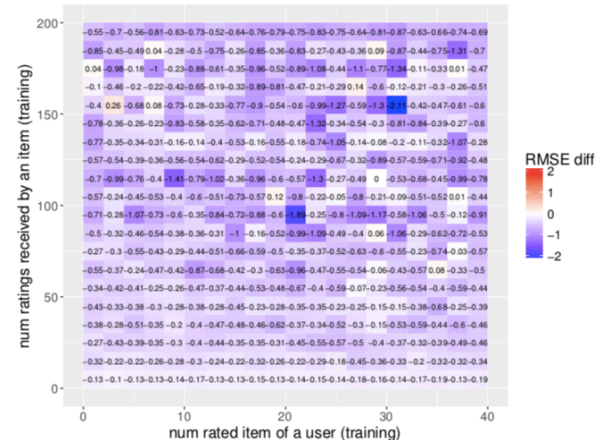
RMSE on Long Tail Items

Dataset	SVD	linear-reg	sqr-reg	log-reg	improve ratio range
Epinions	2.2505	2.1189	2.1186	2.1187	5.85% to 5.86%
MovieLens-100K	1.8866	1.8781	1.8785	1.8808	0.31% to 0.45%
FilmTrust	1.8454	1.7246	1.7204	1.8122	1.80% to 6.77%
Yahoo! Movies	26.1175	25.4348	25.5993	25.5766	1.98% to 2.61%
AMI	3.0426	3.0436	3.0406	3.0448	-0.07% to 0.07%

Dataset	SVD++	linear-reg	sqr-reg	log-reg	improve ratio range
Epinions	2.2295	2.1373	2.1179	2.1186	4.14% to 5.01%
MovieLens-100K	1.8809	1.8797	1.8757	1.8784	0.06% to 0.28%
FilmTrust	1.7523	1.7296	1.7337	1.7180	1.06% to 1.96%
Yahoo! Movies	26.2629	25.5760	25.5558	25.5704	2.62% to 2.69%
AMI	3.1112	3.0448	3.0480	3.0456	2.03% to 2.13%

Dataset	NMF	log-reg	linear-reg	sqr-reg	improve ratio range
Epinions	2.0793	2.0757	2.0753	2.0750	0.17% to 0.21%
MovieLens-100K	2.0788	1.8683	1.8749	1.8724	9.81% to 10.13%
FilmTrust	1.8428	1.6952	1.7026	1.6956	7.61% to 8.01%
Yahoo! Movies	32.8463	25.4302	25.4408	25.4314	22.55% to 22.58%
AMI	2.9942	2.9939	2.9945	2.9950	-0.03% to 0.01%

RMSE diff b2n SVD & SVD+log RDF



16

Discussion

- Determine the **regularization weights** of the latent factors of the item and the user based on their **activeness** is helpful
 - Very simple, can be integrated with many learning-based recommendation models (e.g., SVD, SVD++, NMF, FM, FPMC, Prod2Vec, Behavior2Vec, etc.)
- Common wisdom: larger dataset → smaller regularization weight
 - However, for recommender systems, a larger training dataset may **NOT** necessarily reveal more information of an item or an individual user (esp. long tail items and less active users)

17

If you are interested in details...

- Hung-Hsuan Chen and Pu Chen.
"Differentiating Regularization Weights – a Simple Mechanism to Alleviate Cold Start in Recommender Systems." *ACM Transactions on Knowledge Discovery from Data (TKDD)* 2019.
- GitHub code:
 - <https://github.com/ncu-dart/rdf>

18